

❖ Distributed Deep Learning Architectures

Part 1: Basic Concepts and Mechanisms

1. Data Parallelism vs. Model Parallelism

- **Data Parallelism:**

- **Mechanism:** Every worker holds a **complete copy** of the model and trains on a **different subset of the data**.
- **Synchronization:** Gradients are aggregated (e.g., using **All-Reduce**) after each backward pass.
- **Use Case:** Training **large datasets** with moderate model sizes (e.g., ResNet on ImageNet).

- **Model Parallelism:**

- **Mechanism:** The model itself is **split across multiple devices**. Each device holds only a **part of the model** (e.g., specific layers).
- **Synchronization:** Devices communicate to process the forward/backward pass sequentially across the split model layers.
- **Use Case:** Training **extremely large models** (e.g., GPT-3, BERT-large) that don't fit into a single device's memory.

2. All-Reduce Mechanism

- **Definition:** A **collective communication operation** used in distributed computing that aggregates tensors (typically gradients) across all participating devices and then distributes the resulting aggregated tensor back to every device.

- **Advantages over Parameter Server (PS):**

- **Decentralized:** No central bottleneck, as all nodes communicate peer-to-peer.
- **Scalability:** Offers better scalability and bandwidth utilization due to parallel communication paths.
- **Fault Tolerance:** Lower latency and reduced failure risk since there is no single point of failure (SPOF) like the PS.

3. PyTorch DDP (DistributedDataParallel) Workflow

The DDP module manages synchronous, distributed training across multiple processes (one process per GPU/device):

1. **Setup:** Each process loads the same initial model and a unique **data shard** (managed via `DistributedSampler`).
2. **Local Execution:** Forward and backward passes happen entirely **locally** on each device.
3. **Gradient Sync:** Gradients are averaged across all processes immediately after the backward pass using **All-Reduce**.
4. **Parameter Update:** Each worker uses the aggregated (averaged) gradients to update its **local parameters synchronously**. This ensures all model replicas remain consistent.

4. Core Components of Ray and Object Store Role

- **Core Components:**
 - **Driver:** The main program entry point.
 - **Scheduler:** Manages resources and allocates tasks/actors to worker nodes.
 - **Worker Nodes:** Execute tasks (@ray.remote functions) or actors (@ray.remote classes).
 - **Object Store:** A distributed, in-memory store for immutable objects (data, function return values).
- **Object Store Role:** It facilitates **zero-copy data sharing** between different remote tasks and actors on the same or different nodes. This significantly improves performance by reducing the need for serialization, deserialization, and copying large data structures.

5. Main Limitations of Spark MLlib

Spark MLlib is generally **unsuitable for modern deep learning training** directly because:

- It is primarily focused on **traditional machine learning** algorithms (e.g., logistic regression, decision trees).
- It is **not optimized for GPU computation**, which is essential for deep

learning efficiency.

- Its heavy reliance on the **JVM** introduces inefficiency when interfacing with native Python-based deep learning frameworks (like PyTorch or TensorFlow).
-

6. Role of TorchDistributor in Spark

The `TorchDistributor` package acts as a **bridge** enabling distributed PyTorch (DDP) training to run efficiently on an existing Spark cluster.

- **Core Function:** It launches the required number of PyTorch DDP processes, mapping them directly to **Spark executors**.
 - **Integration:** It manages the environment setup and ensures the communication backend (like NCCL or Gloo) is correctly configured for DDP across the distributed Spark environment.
 - **Benefit:** Allows for a seamless **ETL (Extract, Transform, Load) to training** pipeline within a single, scalable Spark job.
-

Part 2: Framework and Application Questions

7. Ray's Programming Model and Components

Ray's programming model is built on two core abstractions for distributed computing: **Remote Functions** (stateless tasks) and **Actors** (stateful workers).

Abstraction	Decorator	State Management	Purpose
Remote Function	<code>@ray.remote def</code>	Stateless	Defines a function that can be executed
Actor	<code>@ray.remote class</code>	Stateful	Defines a class that can be instantiated as

Example (Conceptual Code Snippets):

```
# Remote Function (Task)
@ray.remote
def square(x):
    return x * x

# Actor (Stateful Worker)
@ray.remote
class Counter:
    def __init__(self):
        self.value = 0
    def increment(self):
        self.value += 1
        return self.value
```

▼ 8. PyTorch DDP Training on Spark using TorchDistributor (Workflow)

The process integrates Spark's data processing with PyTorch's distributed training mechanism.

Schematic Steps:

1. **Spark ETL:** The Spark Driver loads and preprocesses the data into a distributed DataFrame/RDD.
 2. **Job Launch:** The Spark Driver uses **TorchDistributor** to launch the training function onto the cluster.
 3. **DDP Initialization:** Each Spark Executor initializes a process that becomes a **DDP Worker** (assigned a unique rank).
 4. **Data Sharding:** Data shards are distributed/read locally by each DDP Worker.
 5. **Synchronous Training:** DDP Workers train in parallel, with gradients synchronized via **All-Reduce**.
 6. **Checkpoint:** The trained model or checkpoints are saved to distributed storage (e.g., HDFS/S3).
-

9. Ray Train vs. Ray Tune

- **Ray Train:**
 - **Focus:** Handles **distributed model training** itself.
 - **Features:** Provides an API (**Trainer**) for frameworks like PyTorch and TensorFlow, incorporating built-in data parallelism and handling

boilerplate like device management and fault tolerance.

- **Ray Tune:**

- **Focus:** Automates **hyperparameter tuning** and searching for the best model configuration.
 - **Features:** Manages the execution of multiple trials using various search algorithms (e.g., ASHA, HyperBand).
 - **Integration:** Ray Tune is designed to **orchestrate** multiple Ray Train jobs, efficiently distributing them across the cluster to find the optimal training configuration automatically.
-

Part 3: Comprehensive Design and Analysis

11. End-to-End BERT-Large Training Pipeline (Spark + PyTorch + TorchDistributor)

This pipeline leverages the strengths of both systems: Spark for massive-scale data handling and PyTorch/DDP for highly efficient deep learning computation.

Step	Component / Action
1. Data Loading	Spark loads raw text data from distributed storage (S3, HDFS) into a DataFrame.
2. Preprocessing	Text is tokenized (e.g., using a Hugging Face tokenizer) via Spark UDFs or a batch pi
3. Distributed Training	TorchDistributor is initialized on the Spark cluster. The preprocessed data shards ar
4. Synchronization	PyTorch DDP performs forward/backward passes, followed by All-Reduce gradient s
5. Model Saving	The final model or periodic checkpoints are saved back to shared/distributed storage

Conceptual Code Snippet (Spark Launch):

```

from pyspark.sql import SparkSession
from pyspark.ml.torch.distributor import TorchDistributor

# Assume 'preprocessed_data_rdd' is prepared data
def train_fn(iterator):
    # This function runs on each Spark Executor/DDP Worker
    # 1. Setup DDP process group
    # 2. Load model (BERT-large) and optimizer
    # 3. Training loop (read batch from iterator, forward, backward
    pass # ... full DDP training logic

spark = SparkSession.builder.getOrCreate()
distributor = TorchDistributor(num_processes=4, local_mode=False, u

# Run the DDP training across 4 Spark executors with GPUs
distributor.run(train_fn, data=preprocessed_data_rdd)

```

12. Architecture Choice per Scenario

Scenario	Best Architecture	Reason
A: Vision Transformer, multi-GPU, medium data	PyTorch DDP (Native)	Offers the most efficiency
B: Integrated ETL + Training on Spark cluster	Spark + TorchDistributor	Seamlessly leverages Spark's distributed nature
C: Frequent tuning + deployment	Ray Train + Ray Tune + Ray Serve	Provides an end-to-end solution for ML

Part 4: Extended Design & Analysis

13. End-to-End Training Pipeline (DistilBERT Example)

Focusing on a smaller model like DistilBERT highlights the integrated nature of the chosen framework (Spark + PyTorch + TorchDistributor).

- Data Loading & Cleaning (Spark):** Use Spark to read a vast corpus of text data, apply cleaning filters, and deduplicate records across the cluster.
- Tokenization & Tensors (Spark UDFs):** Apply the **DistilBERT tokenizer** as a Spark UDF or use a pre-processing library to convert the text into numerical `input_ids`, `attention_masks`, and `token_type_ids`, storing them as a distributed format suitable for PyTorch loading.
- Distributed Training (TorchDistributor):**
 - The Spark job launches **TorchDistributor**.
 - Each Spark Executor (DDP worker) instantiates the DistilBERT model.

- The workers read their assigned data shards and train synchronously, using **DDP's All-Reduce** for gradient averaging.
4. **Model Saving & Evaluation:** The rank 0 worker saves the final PyTorch state dictionary. Subsequent evaluation runs (either using Spark or another dedicated process) are performed on a hold-out test set.
5. **Deployment:** The saved model is prepared for low-latency inference using a dedicated serving framework like **TorchServe** or **Ray Serve**.
-

14. Architecture Selection (Reanalysis and Final Choice)

Scenario	Recommended Architecture	Final Rationale
A: Vision Transformer, multi-GPU	PyTorch DDP (Native)	Optimal communication
B: ETL + Training on distributed data cluster	Spark + TorchDistributor	Best way to leverage ar
C: Frequent tuning + deployment	Ray Train + Ray Tune + Ray Serve	Provides a unified, flex



Citations / References

- PyTorch Documentation. *Distributed Communication Package - torch.distributed*.
- Ray Documentation. *Ray Core and Ray Train*.
- Apache Spark Documentation. *MLlib and PySpark*.
- Databricks/Spark. *TorchDistributor Documentation*.